

Monitors

Last updated: February 4th 2019, at 12.11pm

Contents

1	Previous methods	1
1.1	Shared variables	1
1.2	Semaphores	2
2	Monitors	2
2.1	The condition data type	2
2.2	Producer/Consumer	2
2.2.1	The Monitor	2
2.2.2	The Producer	3
2.2.3	The Consumer	3
2.3	Problems with monitors	3
2.3.1	Embedded signals	3
2.3.2	Nested monitors	5
3	Monitors and Java	6
3.1	Introduction	6
3.2	Java Code	7
3.2.1	Carpark control	7
3.2.2	Cars arriving	7
3.2.3	Cars leaving	8
3.2.4	The car park	8
3.3	Synchronised statements	9
3.4	Conditions	9

1 Previous methods

1.1 Shared variables

- complex algorithms
- time wasting — busy waits
- primitive (atomic) actions are too weak

Monitors

1.2 Semaphores

- easy to forget to use a semaphore (in complex programmes)
 - forgetting a `poll()` results in lack of mutual exclusion
 - forgetting a `vote()` often leads to deadlock
 - nesting of `poll()`s and `vote()`s is risky
 - cannot test value of semaphore without (possibly) blocking
- Need separate primitives for mutual exclusion and synchronisation.

2 Monitors

A more structured and more user-friendly method of communication.

- generalisation of module concept
- only one process may execute procedures at any one time
- new data type **condition**

2.1 The **condition** data type

Two operations:

- `cond.wait()`
Go to sleep (join the queue for `cond`) and release the monitor
- `cond.signal()`
If the queue for `cond` is not empty, wake the *first* process in the queue.
Otherwise no effect

2.2 Producer/Consumer

Note: this is *not* Java!

2.2.1 The Monitor

```
monitor ProducerConsumer <T> { // this is not Java

    // Conditions for testing
    condition notFull, notEmpty;

    // The buffer
    Buffer<T> buffer;
    ...// see below for producer/consumer code
}
```

2.2.2 The Producer

```
public void producer(T item) {  
    if (buffer.full()) { // is the buffer full?  
        notFull.wait(); // if so, wait (and release the monitor)  
    }  
    buffer.put(item); // add the item to the buffer  
    notEmpty.signal(); // signal that the buffer is not empty  
}
```

2.2.3 The Consumer

```
public T consumer() {  
    if (buffer.empty()) { // is the buffer empty?  
        notEmpty.wait(); // if so, wait (and release the monitor)  
    }  
    T item = buffer.take(); // get the item from the buffer  
    notFull.signal(); // signal that the buffer is not full  
    return item;  
}
```

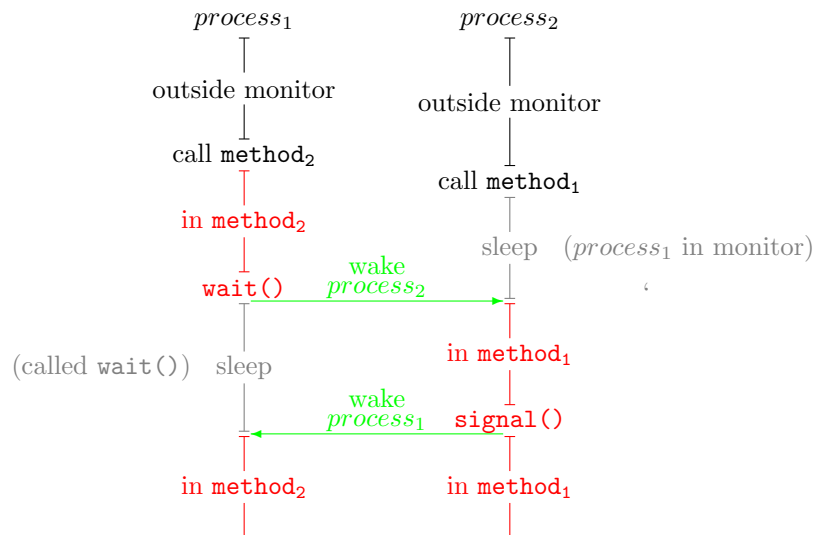
2.3 Problems with monitors

2.3.1 Embedded signals

```
monitor demo {  
    condition cond;  
  
    void method1() {...; cond.signal(); ... }  
  
    void method2() {...; cond.wait(); ... }  
}
```

Can lead to both processes in the monitor simultaneously...

Monitors



Three possible solutions:

- simple: signal and leave
- complex: signal and continue
- even more complex: signal and urgent wait

Simple: Signal and leave

- **signal** *must* be last statement in a monitor
- signalling process will leave monitor immediately after **signal()**
- onus on programmer (or compiler?)

Complex: Signal and continue

- `signal()` wakes a waiting process, if any, and adds it to those contending for the monitor
- signaling process continues until it releases the monitor
- a contending process, if any, obtains the monitor

Signal and urgent wait Process signaling sleeps if necessary.

Two possibilities on `cond.signal()`:

- No processes waiting on `cond`
 - signaller proceeds

Monitors

- Processes waiting on **cond**
 - signaler sleeps
 - wake waiting process
 - wake signaler when this process exits monitor
 - no other process may enter until woken process *and* signaler leave monitor

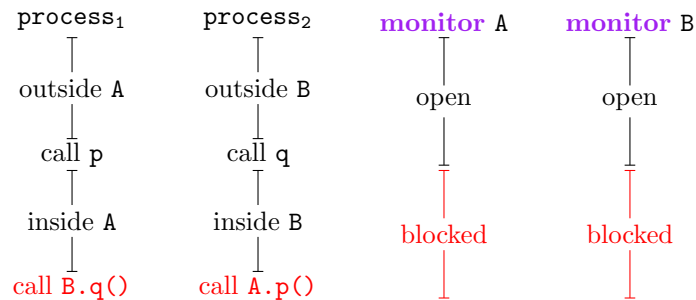
2.3.2 Nested monitors

Mutual calls

```
monitor A {  
    void p() {  
        ...; B.q(); ...  
    }  
}
```

```
monitor B {  
    void q() {  
        ...; A.p(); ...  
    }  
}
```

Can lead to deadlock...



Solution

Assign *priorities* to monitors — no calls to monitors of lower priority allowed.

Blocking waits

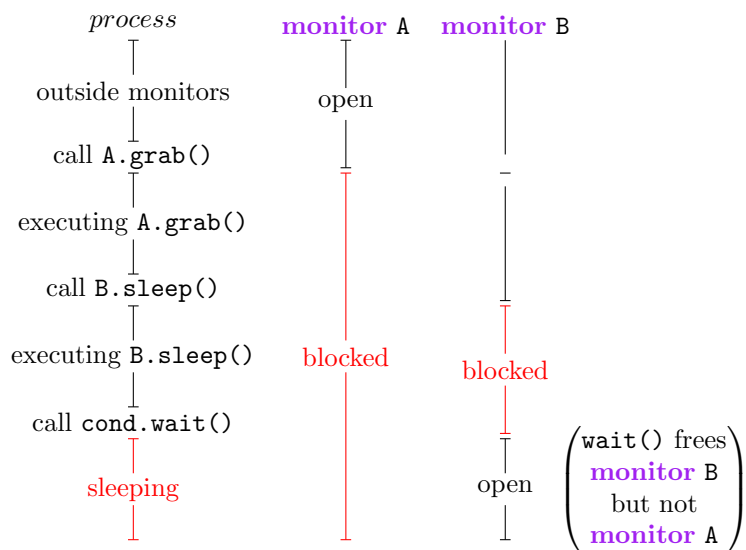
Monitors

```

monitor A {
    void grab() {
        ...; B.sleep(); ...;
    }
    void release() {
        ...; B.wake(); ...
    }
}

monitor B {
    condition cond;
    void sleep() {
        ...; cond.wait(); ...;
    }
    void wake() {
        ...; cond.signal(); ...
    }
}

```



Solution: Process leaves *all* monitors when it blocks — however must then enter multiple monitors when freed (as an atomic action).

3 Monitors and Java

3.1 Introduction

Java — synchronised methods, rather than synchronised classes.

Java keywords:

synchronized method is synchronised (executing thread owns object's monitor)

wait wait until another thread invokes **notify** or **notifyall** (invoking thread must own object's monitor)

Monitors

notify allow a waiting thread, if one exists, to compete for object's monitor
(current thread continues until it relinquishes monitor — signal and continue)

notifyall wake all waiting threads

3.2 Java Code

3.2.1 Carpark control

```
class CarParkControl {

    protected int spaces;
    protected int capacity;

    public CarParkControl(int n) {
        capacity = spaces = n;
    }
    ...// see below for enter() and leave() methods
}

enter()

synchronized void enter() throws InterruptedException {
    if (spaces == 0) {
        wait();
    }
    --spaces;
    notify();
}

leave()

synchronized void leave() throws InterruptedException {
    if (spaces == capacity) {
        wait();
    }
    ++spaces;
    notify();
}
```

3.2.2 Cars arriving

```
class Arrivals extends Thread {
```

Monitors

```
CarParkControl control;

public Arrivals(CarParkControl control) {
    this.control = control;
}

public void run() {
    try {
        while (true) {
            carpark.enter();
        }
    } catch (InterruptedException e) {}
}
```

3.2.3 Cars leaving

```
class Departures extends Thread {

    CarParkControl control;

    public Departures(CarParkControl control) {
        this.control = control;
    }

    public void run() {
        try {
            while (true) {
                carpark.leave();
            }
        } catch (InterruptedException e) {}
    }
}
```

3.2.4 The car park

```
public class CarPark {

    public CarPark(int n) throws InterruptedException {
        CarParkControl control = new CarParkControl(n);
        Arrivals arrivals = new Arrivals(control);
        Departures departures = new Departures(control);
        arrivals.start(); departures.start();
    }
}
```


Monitors

```
        arrivals.join(); departures.join();
    }
}
```

3.3 Synchronised statements

Java also has **synchronized** statements allowing finer grained concurrency. In synchronised statements you need to specify the object whose monitor lock is to be used.

```
private SomeType thing1, thing2;
private Object lock1 = new Object(), lock2 = new Object();

public void access1() {
    synchronized(lock1) {
        // do something to thing1
    }
}
public void access2() {
    synchronized(lock2) {
        // do something to thing2
    }
}
```

3.4 Conditions

Java synchronisation provides you with *one* wait-set per object. Use Conditions and Locks for multiple wait-sets.

```
final Lock lock = new ReentrantLock();
final Condition notFull = lock.newCondition();
final Condition notEmpty = lock.newCondition();

final T[] buffer = new T[...];
final int putIndex = 0; takeIndex = 0; numberOfElements = 0;

public void put(T item) throws InterruptedException {
    lock.lock();
    try {
        while (count == buffer.length) notFull.await();
        buffer[putIndex] = item;
        putIndex = (putIndex+1)%buffer.length;
        numberOfElements++;
        notEmpty.signal();
    }
}
```

Monitors

```
    } finally {  
        lock.unlock();  
    }  
}  
  
public T take() throws InterruptedException {  
    lock.lock();  
    try {  
        while (count == 0) notEmpty.await();  
        T item = buffer[takeIndex];  
        takeIndex = (takeIndex+1)%buffer.length;  
        numberOfElements--;  
        notFull.signal();  
        return item;  
    } finally {  
        lock.unlock();  
    }  
}
```

End of monitors lecture