

Monitors

February 4, 2019

- ▶ Previous methods
- ▶ Monitors
- ▶ Monitors and Java

Previous methods

1: Previous methods

1.1: Shared variables

- ▶ complex algorithms
- ▶ time wasting — busy waits
- ▶ primitive (atomic) actions are too weak

Previous methods

1.2: Semaphores

- ▶ easy to forget to use a semaphore (in complex programmes)
- ▶ forgetting a `poll()` results in lack of mutual exclusion
- ▶ forgetting a `vote()` often leads to deadlock
- ▶ nesting of `poll()`s and `vote()`s is risky
- ▶ cannot test value of semaphore without (possibly) blocking

Need separate primitives for mutual exclusion and synchronisation.

Monitors

2: Monitors

A more structured and more user-friendly method of communication.

- ▶ generalisation of module concept
- ▶ only one process may execute procedures at any one time
- ▶ new data type **condition**

Monitors

2.1: The **condition** data type

Two operations:

- ▶ `cond.wait()`
Go to sleep (join the queue for `cond`) and release the monitor
- ▶ `cond.signal()`
If the queue for `cond` is not empty, wake the *first* process in the queue. Otherwise no effect

Producer/consumer

2.2: Producer/Consumer

2.2.1: The Monitor

```
monitor ProducerConsumer <T> {// this is not Java  
  
    // Conditions for testing  
    condition notFull, notEmpty;  
  
    // The buffer  
    Buffer<T> buffer;  
    ...// see below for producer/consumer code  
}
```

Producer/consumer

2.2.2: The Producer

```
public void producer(T item) {  
    if (buffer.full()) { // is the buffer full?  
        notFull.wait(); // if so, wait (and release the monitor)  
    }  
    buffer.put(item); // add the item to the buffer  
    notEmpty.signal(); // signal that the buffer is not empty  
}
```

Producer/consumer

2.2.3: The Consumer

```
public T consumer() {  
    if (buffer.empty()) { // is the buffer empty?  
        notEmpty.wait(); // if so, wait (and release the monitor)  
    }  
    T item = buffer.take(); // get the item from the buffer  
    notFull.signal(); // signal that the buffer is not full  
    return item;  
}
```


Problems: Embedded signals

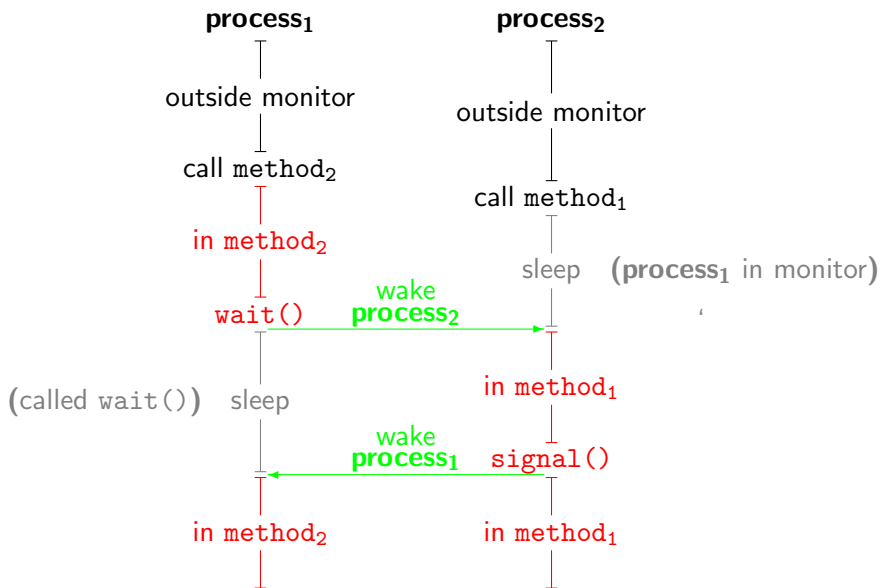
2.3: Problems with monitors

2.3.1: Embedded signals

```
monitor demo {  
    condition cond;  
  
    void method1() {...; cond.signal(); ... }  
  
    void method2() {...; cond.wait(); ... }  
}
```

Can lead to both processes in the monitor simultaneously...

Problems: Embedded signals



Problems: Embedded signals

Three possible solutions:

- ▶ simple: signal and leave
- ▶ complex: signal and continue
- ▶ even more complex: signal and urgent wait

Problems: Embedded signals, Signal and leave

2.3.1 A: Simple: Signal and leave

- ▶ `signal` *must* be last statement in a monitor
- ▶ signalling process will leave monitor immediately after `signal()`
- ▶ onus on programmer (or compiler?)

Problems: Embedded signals, Signal and continue

2.3.1 B: Complex: Signal and continue

- ▶ `signal()` wakes a waiting process, if any, and adds it to those contending for the monitor
- ▶ signaling process continues until it releases the monitor
- ▶ a contending process, if any, obtains the monitor

Problems: Embedded signals, Signal and urgent wait

2.3.1 C: Signal and urgent wait

Process signaling sleeps if necessary.

Two possibilities on `cond.signal()`:

- ▶ No processes waiting on `cond`
 - ▶ signaler proceeds
- ▶ Processes waiting on `cond`
 - ▶ signaler sleeps
 - ▶ wake waiting process
 - ▶ wake signaler when this process exits monitor
 - ▶ no other process may enter until woken process *and* signaler leave monitor

Problems: Nested monitors, Mutual calls

2.3.2: Nested monitors

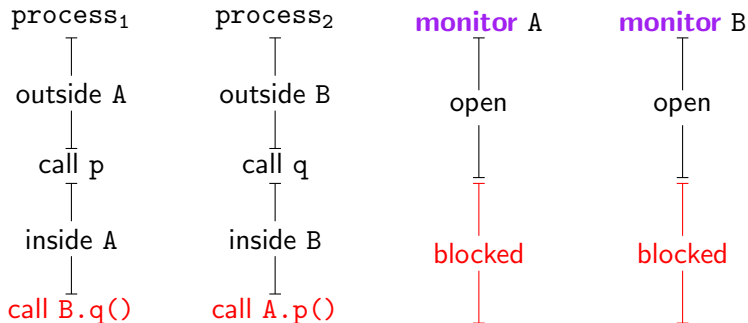
2.3.2 A: Mutual calls

```
monitor A {  
    void p() {  
        ...; B.q(); ...  
    }  
}
```

```
monitor B {  
    void q() {  
        ...; A.p(); ...  
    }  
}
```

Can lead to deadlock...

Problems: Nested monitors, mutual calls



Solution

Assign *priorities* to monitors — no calls to monitors of lower priority allowed.

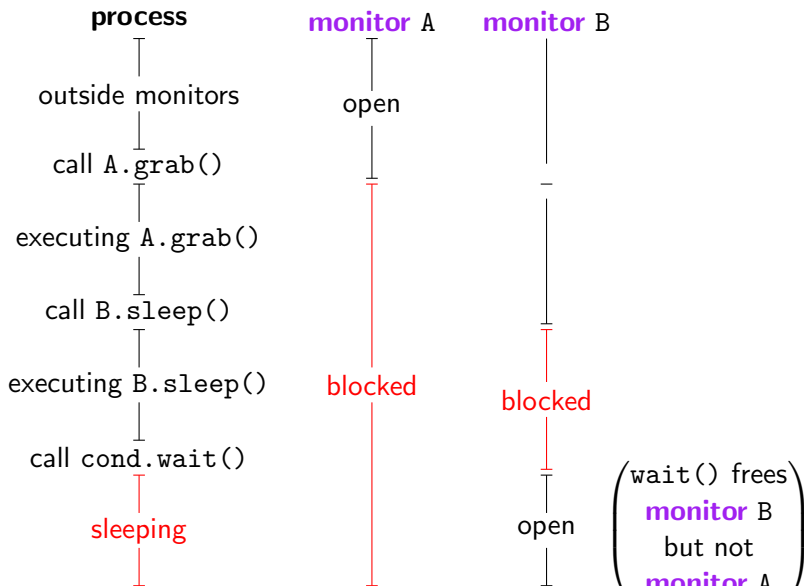
Problems: Nested monitors, Blocking waits

2.3.2 B: Blocking waits

```
monitor A {  
    void grab() {  
        ...; B.sleep(); ...;  
    }  
    void release() {  
        ...; B.wake(); ...  
    }  
}
```

```
monitor B {  
    condition cond;  
    void sleep() {  
        ...; cond.wait(); ...;  
    }  
    void wake() {  
        ...; cond.signal(); ...  
    }  
}
```

Problems: Nested monitors, Blocking waits



Problems: Nested monitors, Blocking waits

Solution:

Process leaves *all* monitors when it blocks — however must then enter multiple monitors when freed (as an atomic action).

3: Monitors and Java

3.1: Introduction

Java — synchronised methods, rather than synchronised classes.

Java keywords:

synchronized method is synchronised (executing thread owns object's monitor)

wait wait until another thread invokes **notify** or **notifyall** (invoking thread must own object's monitor)

notify allow a waiting thread, if one exists, to compete for object's monitor (current thread continues until it relinquishes monitor — signal and continue)

notifyall wake all waiting threads

3.2: Java Code

3.2.1: Carpark control

```
class CarParkControl {  
  
    protected int spaces;  
    protected int capacity;  
  
    public CarParkControl(int n) {  
        capacity = spaces = n;  
    }  
    ...// see below for enter() and leave() methods  
}
```

3.2.1 A: enter()

```
synchronized void enter() throws InterruptedException {  
    if (spaces == 0) {  
        wait();  
    }  
    --spaces;  
    notify();  
}
```

3.2.1 B: leave()

```
synchronized void leave() throws InterruptedException {  
    if (spaces == capacity) {  
        wait();  
    }  
    ++spaces;  
    notify();  
}
```

Java

3.2.2: Cars arriving

```
class Arrivals extends Thread {  
  
    CarParkControl control;  
  
    public Arrivals(CarParkControl control) {  
        this.control = control;  
    }  
  
    public void run() {  
        try {  
            while (true) {  
                carpark.enter();  
            }  
        } catch (InterruptedException e) {}  
    }  
}
```


Java

3.2.3: Cars leaving

```
class Departures extends Thread {  
  
    CarParkControl control;  
  
    public Departures(CarParkControl control) {  
        this.control = control;  
    }  
  
    public void run() {  
        try {  
            while (true) {  
                carpark.leave();  
            }  
        } catch (InterruptedException e) {}  
    }  
}
```

3.2.4: The car park

```
public class CarPark {  
  
    public CarPark(int n) throws InterruptedException {  
        CarParkControl control = new CarParkControl(n);  
        Arrivals arrivals = new Arrivals(control);  
        Departures departures = new Departures(control);  
        arrivals.start(); departures.start();  
        arrivals.join(); departures.join();  
    }  
}
```

3.3: Synchronised statements

Java also has **synchronized** statements. Need to specify the object whose monitor lock is to be used.

Java

```
private SomeType thing1, thing2;
private Object lock1 = new Object(), lock2 = new Object();

public void access1() {
    synchronized(lock1) {
        // do something to thing1
    }
}

public void access2() {
    synchronized(lock2) {
        // do something to thing2
    }
}
```

3.4: Conditions

Java synchronisation provides you with *one* wait-set per object.
Use Conditions and Locks for multiple wait-sets.

```
final Lock lock = new ReentrantLock();  
final Condition notFull = lock.newCondition();  
final Condition notEmpty = lock.newCondition();  
  
final T[] buffer = new T[...];  
final int putIndex = 0; takeIndex = 0; numberOfElements = 0;
```

Java

```
public void put(T item) throws InterruptedException {  
    lock.lock();  
    try {  
        while (count == buffer.length) notFull.await();  
        buffer[putIndex] = item;  
        putIndex = (putIndex+1)%buffer.length;  
        numberOfElements++;  
        notEmpty.signal();  
    } finally {  
        lock.unlock();  
    }  
}
```

Java

```
public T take() throws InterruptedException {  
    lock.lock();  
    try {  
        while (count == 0) notEmpty.await();  
        T item = buffer[takeIndex];  
        takeIndex = (takeIndex+1)%buffer.length;  
        numberOfElements--;  
        notFull.signal();  
        return item;  
    } finally {  
        lock.unlock();  
    }  
}
```

End of monitors lecture